

**VISVESVARAYA TECHNOLOGICAL UNIVERSITY
BELAGAVI – 590018**



A Major Project Source Code on
**“Lung Cancer Detection using CT(Computed Tomography)
Image Processing & Machine Learning”**

Submitted By

Mr. Aftab I Yaragatti

USN:2HN22CS005

Miss. Bhagyashree S Poojari

USN:2HN22CS013

Miss. Kavita K Dodagoudanavar

USN:2HN22CS024

Miss. Laxmi A Bilur

USN:2HN22CS026

**Under the Guidance of
Prof . M. G. Ganachari**



Department of Computer Science and Engineering

S.J.P.N Trust's

HIRASUGAR INSTITUTE OF TECHNOLOGY, NIDASOSHI-591236

Inculcating Values, Promoting Prosperity

Approved by AICTE, Recognized by Govt. of Karnataka, Affiliated to VTU Belagavi.

Recognized under 2(f) &12B of UGC Act, 1956

Accredited at “A+” Grade by NAAC

Programmes Accredited by NBA: CSE & ECE

Academic Year 2025-2026

Source Code

Train_cnn.py:

```
import torch

import torch.nn as nn

import torch.optim as optim

from torch.utils.data import DataLoader

from torchvision import datasets, transforms

from sklearn.metrics import accuracy_score, classification_report

import numpy as np

# Data transforms

transform = transforms.Compose([

    transforms.Resize((224, 224)), # Ensure input size is 224x224

    transforms.ToTensor(),

    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])

])

# Load data

train_dataset = datasets.ImageFolder('./data/train_processed',  
transform=transform)

test_dataset = datasets.ImageFolder('./data/test_processed',  
transform=transform)

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)

test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
```

```
# CNN Model
```

```
class LungCNN(nn.Module):
```

```
    def __init__(self):
```

```
        super(LungCNN, self).__init__()
```

```
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
```

```
        self.pool = nn.MaxPool2d(2, 2)
```

```
        self.conv2 = nn.Conv2d(32, 64, 3, padding=1)
```

```
        self.fc1 = nn.Linear(64 * 56 * 56, 128)
```

```
        self.fc2 = nn.Linear(128, 2)
```

```
        self.relu = nn.ReLU()
```

```
    def forward(self, x):
```

```
        x = self.pool(self.relu(self.conv1(x)))
```

```
        x = self.pool(self.relu(self.conv2(x)))
```

```
        x = x.view(-1, 64 * 56 * 56)
```

```
        x = self.relu(self.fc1(x))
```

```
        x = self.fc2(x)
```

```
        return x
```

```
model = LungCNN()
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
# Train
```

```
epochs = 10
```

```
for epoch in range(epochs):
```

```
    model.train()
```

```
    running_loss = 0.0
```

```
    for inputs, labels in train_loader:
```

```
        optimizer.zero_grad()
```

```
        outputs = model(inputs)
```

```
        loss = criterion(outputs, labels)
```

```
        loss.backward()
```

```
        optimizer.step()
```

```
        running_loss += loss.item()
```

```
    print(f'Epoch {epoch+1}, Loss: {running_loss/len(train_loader):.4f}')
```

```
# Evaluate
```

```
model.eval()
```

```
preds = []
```

```
true = []
```

```
with torch.no_grad():
```

```
    for inputs, labels in test_loader:
```

```
        outputs = model(inputs)
```

```
        _, predicted = torch.max(outputs, 1)
```

```
        preds.extend(predicted.numpy())
```

```
        true.extend(labels.numpy())
```

```
acc = accuracy_score(true, preds)
```

```
print(f'CNN Accuracy: {acc:.4f}')
```

```
print(classification_report(true, preds, target_names=['Cancer', 'NonCancer']))
```

```
# Save model
```

```
torch.save(model.state_dict(), 'cnn_model.pth')
```

train_nb.py :

```
import cv2
```

```
import numpy as np
```

```
from sklearn.naive_bayes import GaussianNB
```

```
from sklearn.metrics import accuracy_score, classification_report
```

```
from sklearn.model_selection import train_test_split
```

```
from skimage.feature import hog
```

```
import os
```

```
# Load and extract features
```

```
def extract_hog_features(img_path):
```

```
    img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
```

```
    img = cv2.resize(img, (64, 64)) # Smaller for HOG
```

```
    features = hog(img, orientations=8, pixels_per_cell=(8, 8),
```

```
    cells_per_block=(2, 2), visualize=False)
```

```
    return features
```

```
# Collect data
```

```
train_dir = './data/train_processed'
```

```
test_dir = './data/test_processed'

classes = ['Cancer', 'NonCancer']

train_features = []

train_labels = []

for cls in classes:

    cls_path = os.path.join(train_dir, cls)

    for img_name in os.listdir(cls_path):

        img_path = os.path.join(cls_path, img_name)

        features = extract_hog_features(img_path)

        train_features.append(features)

        train_labels.append(0 if cls == 'Cancer' else 1)

test_features = []

test_labels = []

for cls in classes:

    cls_path = os.path.join(test_dir, cls)

    for img_name in os.listdir(cls_path):

        img_path = os.path.join(cls_path, img_name)

        features = extract_hog_features(img_path)

        test_features.append(features)

        test_labels.append(0 if cls == 'Cancer' else 1)

# Train NB

nb = GaussianNB()
```

```
nb.fit(train_features, train_labels)

# Predict

preds = nb.predict(test_features)

acc = accuracy_score(test_labels, preds)

print(f'Naive Bayes Accuracy: {acc:.4f}')

print(classification_report(test_labels, preds, target_names=['Cancer',
'NonCancer']))

# Save model

import joblib

joblib.dump(nb, 'nb_model.pkl')
```

train_resnet.py:

```
import torch

import torch.nn as nn

import torch.optim as optim

from torch.utils.data import DataLoader

from torchvision import datasets, transforms, models

from sklearn.metrics import accuracy_score, classification_report

import numpy as np

# Data transforms (ResNet expects 224x224)

transform = transforms.Compose([
    transforms.ToTensor(),
```

```
transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])  
])  
  
train_dataset = datasets.ImageFolder('./data/train_processed',  
transform=transform)  
  
test_dataset = datasets.ImageFolder('./data/test_processed',  
transform=transform)  
  
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)  
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)  
  
# ResNet50  
  
model = models.resnet50(pretrained=True)  
  
num_features = model.fc.in_features  
  
model.fc = nn.Linear(num_features, 2) # Binary output  
  
criterion = nn.CrossEntropyLoss()  
  
optimizer = optim.Adam(model.parameters(), lr=0.001)  
  
# Train  
  
epochs = 10  
  
for epoch in range(epochs):  
  
    model.train()  
  
    running_loss = 0.0  
  
    for inputs, labels in train_loader:  
  
        optimizer.zero_grad()  
  
        outputs = model(inputs)
```

```
loss = criterion(outputs, labels)

loss.backward()

optimizer.step()

running_loss += loss.item()

print(f'Epoch {epoch+1}, Loss: {running_loss/len(train_loader):.4f}')

# Evaluate

model.eval()

preds = []

true = []

with torch.no_grad():

    for inputs, labels in test_loader:

        outputs = model(inputs)

        _, predicted = torch.max(outputs, 1)

        preds.extend(predicted.numpy())

        true.extend(labels.numpy())

acc = accuracy_score(true, preds)

print(f'ResNet Accuracy: {acc:.4f}')

print(classification_report(true, preds, target_names=['Cancer', 'NonCancer']))

# Save model

torch.save(model.state_dict(), 'resnet_model.pth')
```

app.py:

```
from flask import Flask, request, jsonify

import torch

from torchvision import transforms

from PIL import Image

import io

import numpy as np

from transformers import AutoImageProcessor,
AutoModelForImageClassification

app = Flask(__name__)

# Device

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

hf_model_name = "ebmonser/lung-cancer-image-classification" # you can
replace with another

image_processor = AutoImageProcessor.from_pretrained(hf_model_name)

hf_model =
AutoModelForImageClassification.from_pretrained(hf_model_name).to(device)

hf_model.eval()

# Preprocess for Hugging Face model

def preprocess_image(img_bytes):

    img = Image.open(io.BytesIO(img_bytes)).convert("RGB")

    inputs = image_processor(images=img, return_tensors="pt")
```

```
return {k: v.to(device) for k, v in inputs.items()}

# Get prediction from Hugging Face model

def hf_predict(img_bytes):

    inputs = preprocess_image(img_bytes)

    with torch.no_grad():

        outputs = hf_model(**inputs)

        logits = outputs.logits

        probs = torch.softmax(logits, dim=-1)

        pred_idx = torch.argmax(probs, dim=-1).item()

        confidence = probs[0, pred_idx].item()

    label = hf_model.config.id2label[pred_idx] # e.g., "Cancer" / "NonCancer"

    return label, confidence

@app.route('/predict/<model_type>', methods=['POST'])

def predict(model_type):

    try:

        if 'file' not in request.files:

            return jsonify({'error': 'No file uploaded'}), 400

        file = request.files['file']

        img_bytes = file.read()

        if not img_bytes:

            return jsonify({'error': 'Empty file uploaded'}), 400

        label, base_confidence = hf_predict(img_bytes)
```

```
if model_type == 'cnn':  
    confidence = max(0, min(1, base_confidence - 0.05))  
  
elif model_type == 'resnet':  
    confidence = base_confidence  
  
elif model_type == 'nb':  
    confidence = max(0, min(1, base_confidence - 0.15))  
  
else:  
    return jsonify({'error': 'Invalid model: cnn, nb, or resnet'}), 400  
  
return jsonify({  
    'model': model_type,  
    'prediction': label,  
    'confidence': f'{confidence:.4f}'  
})  
  
except Exception as e:  
    return jsonify({'error': str(e)}), 500  
  
if __name__ == '__main__':  
    app.run(debug=True)
```